

Capstone_Project_tf_course_one

November 22, 2020

1 Capstone Project

1.1 Image classifier for the SVHN dataset

1.1.1 Instructions

In this notebook, you will create a neural network that classifies real-world images digits. You will use concepts from throughout this course in building, training, testing, validating and saving your Tensorflow classifier model.

This project is peer-assessed. Within this notebook you will find instructions in each section for how to complete the project. Pay close attention to the instructions as the peer review will be carried out according to a grading rubric that checks key parts of the project instructions. Feel free to add extra cells into the notebook as required.

1.1.2 How to submit

When you have completed the Capstone project notebook, you will submit a pdf of the notebook for peer review. First ensure that the notebook has been fully executed from beginning to end, and all of the cell outputs are visible. This is important, as the grading rubric depends on the reviewer being able to view the outputs of your notebook. Save the notebook as a pdf (you could download the notebook with File -> Download .ipynb, open the notebook locally, and then File -> Download as -> PDF via LaTeX), and then submit this pdf for review.

1.1.3 Let's get started!

We'll start by running some imports, and loading the dataset. For this project you are free to make further imports throughout the notebook as you wish.

```
[1]: import tensorflow as tf
    from scipy.io import loadmat
```

For the capstone project, you will use the [SVHN dataset](#). This is an image dataset of over 600,000 digit images in all, and is a harder dataset than MNIST as the numbers appear in the context of natural scene images. SVHN is obtained from house numbers in Google Street View images.

- Y. Netzer, T. Wang, A. Coates, A. Bissacco, B. Wu and A. Y. Ng. "Reading Digits in Natural Images with Unsupervised Feature Learning". NIPS Workshop on Deep Learning and Unsupervised Feature Learning, 2011.

The train and test datasets required for this project can be downloaded from [here](#) and [here](#). Once unzipped, you will have two files: `train_32x32.mat` and `test_32x32.mat`. You should store these files in Drive for use in this Colab notebook.

Your goal is to develop an end-to-end workflow for building, training, validating, evaluating and saving a neural network that classifies a real-world image into one of ten classes.

```
[2]: # Run this cell to connect to your Drive folder
```

```
from google.colab import drive
drive.mount('/content/gdrive')
```

Mounted at /content/gdrive

```
[3]: # Load the dataset from your Drive folder
```

```
train = loadmat('/content/gdrive/MyDrive/Colab Notebooks/train_32x32.mat')
test = loadmat('/content/gdrive/MyDrive/Colab Notebooks/test_32x32.mat')
```

Both `train` and `test` are dictionaries with keys `X` and `y` for the input images and labels respectively.

1.2 1. Inspect and preprocess the dataset

- Extract the training and testing images and labels separately from the `train` and `test` dictionaries loaded for you.
- Select a random sample of images and corresponding labels from the dataset (at least 10), and display them in a figure.
- Convert the training and test images to grayscale by taking the average across all colour channels for each pixel. *Hint: retain the channel dimension, which will now have size 1.*
- Select a random sample of the grayscale images and corresponding labels from the dataset (at least 10), and display them in a figure.

```
[4]: import numpy
train_data=train['X']
train_data=numpy.moveaxis(train_data,3,0)
train_data=train_data/255.
train_labels=train['y']
test_data=test['X']
test_data=numpy.moveaxis(test_data,3,0)
test_data=test_data/255.
test_labels=test['y']
print(train_data.shape)
print(test_data.shape)
print(train_labels.shape)
```

```
(73257, 32, 32, 3)
(26032, 32, 32, 3)
(73257, 1)
```


```
[5]: import numpy as np
import matplotlib.pyplot as plt
num_train=train_data.shape[0]
np.random.seed(0)
random_idx=np.random.choice(num_train,10)
fig,axes=plt.subplots(1,10,figsize=(10,1))
for i in range(10):
    axes[i].set_axis_off()
    axes[i].imshow(train_data[random_idx[i]])
    axes[i].text(2., -1.5, f'Digit {train_labels[random_idx[i],,]}')
```



```
[6]: train_data_gray=np.mean(train_data,axis=3)
test_data_gray=np.mean(test_data,axis=3)
train_data_gray=train_data_gray[...,np.newaxis]
test_data_gray=test_data_gray[...,np.newaxis]
print(train_data_gray.shape)
```

(73257, 32, 32, 1)

```
[7]: np.random.seed(0)
random_idx_1=np.random.choice(num_train,10)
fig,axes=plt.subplots(1,10,figsize=(10,1))
for i in range(10):
    axes[i].set_axis_off()
    axes[i].imshow(np.squeeze(train_data_gray[random_idx_1[i]]))
    axes[i].text(2., -1.5, f'Digit {train_labels[random_idx_1[i],,]}')
```



1.3 2. MLP neural network classifier

- Build an MLP classifier model using the Sequential API. Your model should use only Flatten and Dense layers, with the final layer having a 10-way softmax output.

- You should design and build the model yourself. Feel free to experiment with different MLP architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 4 or 5 layers.*
- Print out the model summary (using the `summary()` method)
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a `ModelCheckpoint` callback.
- As a guide, you should aim to achieve a final categorical cross entropy training loss of less than 1.0 (the validation loss might be higher).
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
[8]: from tensorflow.keras.models import Sequential
      from tensorflow.keras.layers import Dense, Conv2D,
      ↪MaxPooling2D, Flatten, BatchNormalization, Dropout
```

```
[9]: def get_new_model(input_shape):
      model=Sequential([
          Flatten(input_shape=(32, 32)),
          Dense(512, activation='tanh', name='Dense_1'),
          Dense(128, activation='tanh', name='Dense_2'),
          Dense(128, activation='tanh', name='Dense_3'),
          Dense(10, activation='softmax', name='final'),
      ])
      return model
```

```
[10]: model=get_new_model(train_data_gray[0].shape)
      model.summary()
```

Model: "sequential"

Layer (type)	Output Shape	Param #
flatten (Flatten)	(None, 1024)	0
Dense_1 (Dense)	(None, 512)	524800
Dense_2 (Dense)	(None, 128)	65664
Dense_3 (Dense)	(None, 128)	16512
final (Dense)	(None, 10)	1290
Total params: 608,266		
Trainable params: 608,266		
Non-trainable params: 0		

```

[11]: opt=tf.keras.optimizers.Adam(learning_rate=0.001)
      model.compile(optimizer=opt,loss='categorical_crossentropy',metrics=['accuracy'])

[12]: from tensorflow.keras.callbacks import ModelCheckpoint
      #early_stopping=tf.keras.callbacks.EarlyStopping(patience=3)
      epoch_accuracy_callback=tf.keras.callbacks.LambdaCallback(
          on_epoch_end=lambda epoch,logs:print('\n after_
          ↳epoch{0},the accuracy is {0:7.2f}'.format(epoch,logs['accuracy'])))
      checkpoint_best_only_path='checkpoint_best_only/checkpoint'
      checkpoint_best_only=ModelCheckpoint(filepath=checkpoint_best_only_path,
          save_weights_only=True,
          save_freq='epoch',
          save_best_only=True,
          monitor='val_accuracy',
          verbose=1)

[13]: train_labels=[i if i!=10 else 0 for i in train_labels ]
      train_labels=tf.keras.utils.to_categorical(np.array(train_labels))

[14]: history=model.fit(train_data_gray,
          train_labels,
          validation_split=0.2,
          epochs=30,
          batch_size=64,
          verbose=1,
          callbacks=[epoch_accuracy_callback,checkpoint_best_only])

```

Epoch 1/30

915/916 [=====>.] - ETA: 0s - loss: 2.2277 - accuracy: 0.1881

after epoch0,the accuracy is 0.19

Epoch 00001: val_accuracy improved from -inf to 0.24939, saving model to checkpoint_best_only/checkpoint

916/916 [=====] - 7s 8ms/step - loss: 2.2275 - accuracy: 0.1882 - val_loss: 2.1140 - val_accuracy: 0.2494

Epoch 2/30

912/916 [=====>.] - ETA: 0s - loss: 1.9169 - accuracy: 0.3169

after epoch1,the accuracy is 0.32

Epoch 00002: val_accuracy improved from 0.24939 to 0.42540, saving model to checkpoint_best_only/checkpoint

916/916 [=====] - 7s 8ms/step - loss: 1.9160 - accuracy: 0.3172 - val_loss: 1.6446 - val_accuracy: 0.4254

Epoch 3/30

916/916 [=====] - ETA: 0s - loss: 1.5647 - accuracy: 0.4561

after epoch2,the accuracy is 0.46

Epoch 00003: val_accuracy improved from 0.42540 to 0.48990, saving model to
checkpoint_best_only/checkpoint
916/916 [=====] - 7s 8ms/step - loss: 1.5647 -
accuracy: 0.4561 - val_loss: 1.4971 - val_accuracy: 0.4899
Epoch 4/30
910/916 [=====>.] - ETA: 0s - loss: 1.4026 - accuracy:
0.5268
after epoch3,the accuracy is 0.53

Epoch 00004: val_accuracy improved from 0.48990 to 0.57501, saving model to
checkpoint_best_only/checkpoint
916/916 [=====] - 7s 8ms/step - loss: 1.4015 -
accuracy: 0.5272 - val_loss: 1.3121 - val_accuracy: 0.5750
Epoch 5/30
912/916 [=====>.] - ETA: 0s - loss: 1.2684 - accuracy:
0.5859
after epoch4,the accuracy is 0.59

Epoch 00005: val_accuracy improved from 0.57501 to 0.60579, saving model to
checkpoint_best_only/checkpoint
916/916 [=====] - 7s 8ms/step - loss: 1.2682 -
accuracy: 0.5860 - val_loss: 1.2273 - val_accuracy: 0.6058
Epoch 6/30
911/916 [=====>.] - ETA: 0s - loss: 1.1834 - accuracy:
0.6227
after epoch5,the accuracy is 0.62

Epoch 00006: val_accuracy improved from 0.60579 to 0.62851, saving model to
checkpoint_best_only/checkpoint
916/916 [=====] - 7s 8ms/step - loss: 1.1834 -
accuracy: 0.6228 - val_loss: 1.1753 - val_accuracy: 0.6285
Epoch 7/30
910/916 [=====>.] - ETA: 0s - loss: 1.1315 - accuracy:
0.6414
after epoch6,the accuracy is 0.64

Epoch 00007: val_accuracy did not improve from 0.62851
916/916 [=====] - 7s 8ms/step - loss: 1.1315 -
accuracy: 0.6415 - val_loss: 1.2207 - val_accuracy: 0.6103
Epoch 8/30
913/916 [=====>.] - ETA: 0s - loss: 1.0918 - accuracy:
0.6555
after epoch7,the accuracy is 0.66

Epoch 00008: val_accuracy improved from 0.62851 to 0.65192, saving model to
checkpoint_best_only/checkpoint

916/916 [=====] - 7s 8ms/step - loss: 1.0921 - accuracy: 0.6553 - val_loss: 1.1086 - val_accuracy: 0.6519
Epoch 9/30
910/916 [=====>.] - ETA: 0s - loss: 1.0583 - accuracy: 0.6660
after epoch8,the accuracy is 0.67

Epoch 00009: val_accuracy improved from 0.65192 to 0.66510, saving model to checkpoint_best_only/checkpoint
916/916 [=====] - 7s 8ms/step - loss: 1.0578 - accuracy: 0.6661 - val_loss: 1.0705 - val_accuracy: 0.6651
Epoch 10/30
912/916 [=====>.] - ETA: 0s - loss: 1.0351 - accuracy: 0.6755
after epoch9,the accuracy is 0.68

Epoch 00010: val_accuracy improved from 0.66510 to 0.68271, saving model to checkpoint_best_only/checkpoint
916/916 [=====] - 7s 8ms/step - loss: 1.0353 - accuracy: 0.6755 - val_loss: 1.0166 - val_accuracy: 0.6827
Epoch 11/30
911/916 [=====>.] - ETA: 0s - loss: 1.0066 - accuracy: 0.6841
after epoch10,the accuracy is 0.68

Epoch 00011: val_accuracy did not improve from 0.68271
916/916 [=====] - 7s 8ms/step - loss: 1.0070 - accuracy: 0.6838 - val_loss: 1.0908 - val_accuracy: 0.6564
Epoch 12/30
915/916 [=====>.] - ETA: 0s - loss: 0.9894 - accuracy: 0.6910
after epoch11,the accuracy is 0.69

Epoch 00012: val_accuracy did not improve from 0.68271
916/916 [=====] - 7s 8ms/step - loss: 0.9893 - accuracy: 0.6911 - val_loss: 1.0090 - val_accuracy: 0.6807
Epoch 13/30
916/916 [=====] - ETA: 0s - loss: 0.9713 - accuracy: 0.6959
after epoch12,the accuracy is 0.70

Epoch 00013: val_accuracy improved from 0.68271 to 0.68851, saving model to checkpoint_best_only/checkpoint
916/916 [=====] - 7s 8ms/step - loss: 0.9713 - accuracy: 0.6959 - val_loss: 0.9828 - val_accuracy: 0.6885
Epoch 14/30
911/916 [=====>.] - ETA: 0s - loss: 0.9560 - accuracy: 0.7019

after epoch13,the accuracy is 0.70

Epoch 00014: val_accuracy improved from 0.68851 to 0.69608, saving model to checkpoint_best_only/checkpoint
 916/916 [=====] - 7s 8ms/step - loss: 0.9554 - accuracy: 0.7023 - val_loss: 0.9839 - val_accuracy: 0.6961
 Epoch 15/30
 913/916 [=====>.] - ETA: 0s - loss: 0.9404 - accuracy: 0.7063
 after epoch14,the accuracy is 0.71

Epoch 00015: val_accuracy improved from 0.69608 to 0.70141, saving model to checkpoint_best_only/checkpoint
 916/916 [=====] - 10s 11ms/step - loss: 0.9401 - accuracy: 0.7063 - val_loss: 0.9539 - val_accuracy: 0.7014
 Epoch 16/30
 916/916 [=====] - ETA: 0s - loss: 0.9236 - accuracy: 0.7135
 after epoch15,the accuracy is 0.71

Epoch 00016: val_accuracy improved from 0.70141 to 0.70577, saving model to checkpoint_best_only/checkpoint
 916/916 [=====] - 8s 9ms/step - loss: 0.9236 - accuracy: 0.7135 - val_loss: 0.9561 - val_accuracy: 0.7058
 Epoch 17/30
 911/916 [=====>.] - ETA: 0s - loss: 0.9104 - accuracy: 0.7177
 after epoch16,the accuracy is 0.72

Epoch 00017: val_accuracy did not improve from 0.70577
 916/916 [=====] - 7s 8ms/step - loss: 0.9112 - accuracy: 0.7175 - val_loss: 0.9760 - val_accuracy: 0.6953
 Epoch 18/30
 916/916 [=====] - ETA: 0s - loss: 0.9041 - accuracy: 0.7184
 after epoch17,the accuracy is 0.72

Epoch 00018: val_accuracy did not improve from 0.70577
 916/916 [=====] - 7s 8ms/step - loss: 0.9041 - accuracy: 0.7184 - val_loss: 1.0041 - val_accuracy: 0.6857
 Epoch 19/30
 910/916 [=====>.] - ETA: 0s - loss: 0.8932 - accuracy: 0.7227
 after epoch18,the accuracy is 0.72

Epoch 00019: val_accuracy did not improve from 0.70577
 916/916 [=====] - 7s 8ms/step - loss: 0.8931 - accuracy: 0.7229 - val_loss: 0.9507 - val_accuracy: 0.7006

Epoch 20/30
915/916 [=====>.] - ETA: 0s - loss: 0.8870 - accuracy: 0.7248
after epoch19,the accuracy is 0.72

Epoch 00020: val_accuracy improved from 0.70577 to 0.70782, saving model to checkpoint_best_only/checkpoint
916/916 [=====] - 7s 8ms/step - loss: 0.8871 - accuracy: 0.7248 - val_loss: 0.9393 - val_accuracy: 0.7078

Epoch 21/30
916/916 [=====] - ETA: 0s - loss: 0.8680 - accuracy: 0.7319
after epoch20,the accuracy is 0.73

Epoch 00021: val_accuracy did not improve from 0.70782
916/916 [=====] - 8s 8ms/step - loss: 0.8680 - accuracy: 0.7319 - val_loss: 0.9589 - val_accuracy: 0.6983

Epoch 22/30
916/916 [=====] - ETA: 0s - loss: 0.8616 - accuracy: 0.7331
after epoch21,the accuracy is 0.73

Epoch 00022: val_accuracy improved from 0.70782 to 0.71963, saving model to checkpoint_best_only/checkpoint
916/916 [=====] - 8s 8ms/step - loss: 0.8616 - accuracy: 0.7331 - val_loss: 0.9047 - val_accuracy: 0.7196

Epoch 23/30
914/916 [=====>.] - ETA: 0s - loss: 0.8619 - accuracy: 0.7317
after epoch22,the accuracy is 0.73

Epoch 00023: val_accuracy did not improve from 0.71963
916/916 [=====] - 7s 8ms/step - loss: 0.8620 - accuracy: 0.7316 - val_loss: 0.9474 - val_accuracy: 0.7040

Epoch 24/30
912/916 [=====>.] - ETA: 0s - loss: 0.8406 - accuracy: 0.7390
after epoch23,the accuracy is 0.74

Epoch 00024: val_accuracy improved from 0.71963 to 0.73567, saving model to checkpoint_best_only/checkpoint
916/916 [=====] - 8s 8ms/step - loss: 0.8406 - accuracy: 0.7391 - val_loss: 0.8724 - val_accuracy: 0.7357

Epoch 25/30
912/916 [=====>.] - ETA: 0s - loss: 0.8469 - accuracy: 0.7366
after epoch24,the accuracy is 0.74

Epoch 00025: val_accuracy did not improve from 0.73567
 916/916 [=====] - 7s 8ms/step - loss: 0.8465 - accuracy: 0.7368 - val_loss: 0.8678 - val_accuracy: 0.7297
 Epoch 26/30
 914/916 [=====>.] - ETA: 0s - loss: 0.8385 - accuracy: 0.7392
 after epoch25,the accuracy is 0.74

Epoch 00026: val_accuracy did not improve from 0.73567
 916/916 [=====] - 7s 8ms/step - loss: 0.8385 - accuracy: 0.7393 - val_loss: 0.8684 - val_accuracy: 0.7347
 Epoch 27/30
 910/916 [=====>.] - ETA: 0s - loss: 0.8325 - accuracy: 0.7411
 after epoch26,the accuracy is 0.74

Epoch 00027: val_accuracy did not improve from 0.73567
 916/916 [=====] - 7s 8ms/step - loss: 0.8321 - accuracy: 0.7413 - val_loss: 0.9103 - val_accuracy: 0.7209
 Epoch 28/30
 913/916 [=====>.] - ETA: 0s - loss: 0.8203 - accuracy: 0.7460
 after epoch27,the accuracy is 0.75

Epoch 00028: val_accuracy did not improve from 0.73567
 916/916 [=====] - 7s 8ms/step - loss: 0.8206 - accuracy: 0.7458 - val_loss: 0.8836 - val_accuracy: 0.7256
 Epoch 29/30
 909/916 [=====>.] - ETA: 0s - loss: 0.8209 - accuracy: 0.7449
 after epoch28,the accuracy is 0.74

Epoch 00029: val_accuracy did not improve from 0.73567
 916/916 [=====] - 7s 8ms/step - loss: 0.8214 - accuracy: 0.7449 - val_loss: 0.8965 - val_accuracy: 0.7232
 Epoch 30/30
 911/916 [=====>.] - ETA: 0s - loss: 0.8208 - accuracy: 0.7445
 after epoch29,the accuracy is 0.74

Epoch 00030: val_accuracy did not improve from 0.73567
 916/916 [=====] - 8s 8ms/step - loss: 0.8207 - accuracy: 0.7446 - val_loss: 0.8641 - val_accuracy: 0.7307

```
[15]: import matplotlib.pyplot as plt
import pandas as pd
fig=plt.figure(figsize=(12,5))
```

```

fig.add_subplot(121)


plt.plot(history.history['loss'])
plt.plot(history.history['val_loss'])
plt.title('loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

fig.add_subplot(122)

plt.plot(history.history['accuracy'])
plt.plot(history.history['val_accuracy'])
plt.title('accuracy vs. epochs')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

plt.show()

```



```

[16]: test_labels=[i if i!=10 else 0 for i in test_labels ]
test_labels=tf.keras.utils.to_categorical(np.array(test_labels))
model.evaluate(test_data_gray,test_labels)

```

814/814 [=====] - 2s 2ms/step - loss: 0.9387 - accuracy: 0.7120

[16]: [0.9386504292488098, 0.7120466828346252]

1.4 3. CNN neural network classifier

- Build a CNN classifier model using the Sequential API. Your model should use the Conv2D, MaxPool2D, BatchNormalization, Flatten, Dense and Dropout layers. The final layer should again have a 10-way softmax output.
- You should design and build the model yourself. Feel free to experiment with different CNN architectures. *Hint: to achieve a reasonable accuracy you won't need to use more than 2 or 3 convolutional layers and 2 fully connected layers.*
- The CNN model should use fewer trainable parameters than your MLP model.
- Compile and train the model (we recommend a maximum of 30 epochs), making use of both training and validation sets during the training run.
- Your model should track at least one appropriate metric, and use at least two callbacks during training, one of which should be a ModelCheckpoint callback.
- You should aim to beat the MLP model performance with fewer parameters!
- Plot the learning curves for loss vs epoch and accuracy vs epoch for both training and validation sets.
- Compute and display the loss and accuracy of the trained model on the test set.

```
[17]: model_cnn=Sequential([  
    ↳  
    ↳Conv2D(16,kernel_size=3,activation='tanh',input_shape=(32,32,1)),  
        MaxPooling2D(pool_size=3),  
        Dense(64,activation='tanh'),  
        Flatten(),  
        Dense(10,activation='softmax')  
])
```

```
[18]: model_cnn.summary()
```

Model: "sequential_1"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 30, 30, 16)	160
max_pooling2d (MaxPooling2D)	(None, 10, 10, 16)	0
dense (Dense)	(None, 10, 10, 64)	1088
flatten_1 (Flatten)	(None, 6400)	0
dense_1 (Dense)	(None, 10)	64010

=====
Total params: 65,258
Trainable params: 65,258
Non-trainable params: 0
=====

```
[19]: opt=tf.keras.optimizers.Adam(learning_rate=0.001)
model_cnn.
    ↳ compile(optimizer=opt,loss='categorical_crossentropy',metrics=['accuracy'])
epoch_accuracy_callback_cnn=tf.keras.callbacks.LambdaCallback(
    on_epoch_end=lambda epoch,logs:print('\n after_
    ↳ epoch{},the accuracy is {:.2f}'.format(epoch,logs['accuracy'])))
checkpoint_best_only_cnn_path='checkpoint_best_only_cnn/checkpoint'
checkpoint_best_only_cnn=ModelCheckpoint(filepath=checkpoint_best_only_cnn_path,
    save_weights_only=True,
    save_freq='epoch',
    save_best_only=True,
    monitor='val_accuracy',
    verbose=1)

[20]: history_cnn=model_cnn.fit(train_data_gray,
    train_labels,
    validation_split=0.2,
    epochs=30,
    batch_size=64,
    verbose=1,
    callbacks=[epoch_accuracy_callback_cnn,checkpoint_best_only_cnn])
```

```
Epoch 1/30
915/916 [=====>.] - ETA: 0s - loss: 1.3787 - accuracy:
0.5619
after epoch0,the accuracy is    0.56
```

```
Epoch 00001: val_accuracy improved from -inf to 0.73956, saving model to
checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 28ms/step - loss: 1.3785 -
accuracy: 0.5620 - val_loss: 0.9244 - val_accuracy: 0.7396
```

```
Epoch 2/30
915/916 [=====>.] - ETA: 0s - loss: 0.7919 - accuracy:
0.7742
after epoch1,the accuracy is    0.77
```

```
Epoch 00002: val_accuracy improved from 0.73956 to 0.79928, saving model to
checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 28ms/step - loss: 0.7920 -
accuracy: 0.7742 - val_loss: 0.7415 - val_accuracy: 0.7993
```

```
Epoch 3/30
915/916 [=====>.] - ETA: 0s - loss: 0.6879 - accuracy:
0.8076
after epoch2,the accuracy is    0.81
```

```
Epoch 00003: val_accuracy improved from 0.79928 to 0.80685, saving model to
checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 27ms/step - loss: 0.6882 -
```

accuracy: 0.8076 - val_loss: 0.7129 - val_accuracy: 0.8069
Epoch 4/30
915/916 [=====>.] - ETA: 0s - loss: 0.6601 - accuracy: 0.8154
after epoch3,the accuracy is 0.82

Epoch 00004: val_accuracy improved from 0.80685 to 0.80699, saving model to checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 27ms/step - loss: 0.6599 - accuracy: 0.8154 - val_loss: 0.7055 - val_accuracy: 0.8070
Epoch 5/30
916/916 [=====] - ETA: 0s - loss: 0.6415 - accuracy: 0.8214
after epoch4,the accuracy is 0.82

Epoch 00005: val_accuracy improved from 0.80699 to 0.81402, saving model to checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 27ms/step - loss: 0.6415 - accuracy: 0.8214 - val_loss: 0.6915 - val_accuracy: 0.8140
Epoch 6/30
916/916 [=====] - ETA: 0s - loss: 0.6297 - accuracy: 0.8242
after epoch5,the accuracy is 0.82

Epoch 00006: val_accuracy improved from 0.81402 to 0.81600, saving model to checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 27ms/step - loss: 0.6297 - accuracy: 0.8242 - val_loss: 0.6887 - val_accuracy: 0.8160
Epoch 7/30
914/916 [=====>.] - ETA: 0s - loss: 0.6198 - accuracy: 0.8272
after epoch6,the accuracy is 0.83

Epoch 00007: val_accuracy improved from 0.81600 to 0.81743, saving model to checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 26s 28ms/step - loss: 0.6201 - accuracy: 0.8272 - val_loss: 0.6863 - val_accuracy: 0.8174
Epoch 8/30
914/916 [=====>.] - ETA: 0s - loss: 0.6114 - accuracy: 0.8283
after epoch7,the accuracy is 0.83

Epoch 00008: val_accuracy improved from 0.81743 to 0.82064, saving model to checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 27ms/step - loss: 0.6111 - accuracy: 0.8284 - val_loss: 0.6753 - val_accuracy: 0.8206
Epoch 9/30
914/916 [=====>.] - ETA: 0s - loss: 0.6034 - accuracy:

0.8334
after epoch8,the accuracy is 0.83

Epoch 00009: val_accuracy did not improve from 0.82064
916/916 [=====] - 25s 28ms/step - loss: 0.6036 -
accuracy: 0.8334 - val_loss: 0.6874 - val_accuracy: 0.8168
Epoch 10/30
914/916 [=====>.] - ETA: 0s - loss: 0.5937 - accuracy:
0.8345
after epoch9,the accuracy is 0.83

Epoch 00010: val_accuracy improved from 0.82064 to 0.82098, saving model to
checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 28ms/step - loss: 0.5936 -
accuracy: 0.8346 - val_loss: 0.6734 - val_accuracy: 0.8210
Epoch 11/30
915/916 [=====>.] - ETA: 0s - loss: 0.5863 - accuracy:
0.8371
after epoch10,the accuracy is 0.84

Epoch 00011: val_accuracy did not improve from 0.82098
916/916 [=====] - 25s 27ms/step - loss: 0.5863 -
accuracy: 0.8371 - val_loss: 0.6799 - val_accuracy: 0.8181
Epoch 12/30
916/916 [=====] - ETA: 0s - loss: 0.5815 - accuracy:
0.8378
after epoch11,the accuracy is 0.84

Epoch 00012: val_accuracy improved from 0.82098 to 0.82159, saving model to
checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 27ms/step - loss: 0.5815 -
accuracy: 0.8378 - val_loss: 0.6682 - val_accuracy: 0.8216
Epoch 13/30
914/916 [=====>.] - ETA: 0s - loss: 0.5744 - accuracy:
0.8379
after epoch12,the accuracy is 0.84

Epoch 00013: val_accuracy improved from 0.82159 to 0.82651, saving model to
checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 25s 28ms/step - loss: 0.5744 -
accuracy: 0.8379 - val_loss: 0.6607 - val_accuracy: 0.8265
Epoch 14/30
915/916 [=====>.] - ETA: 0s - loss: 0.5687 - accuracy:
0.8417
after epoch13,the accuracy is 0.84

Epoch 00014: val_accuracy did not improve from 0.82651
916/916 [=====] - 26s 28ms/step - loss: 0.5686 -

```

accuracy: 0.8417 - val_loss: 0.6660 - val_accuracy: 0.8245
Epoch 15/30
915/916 [=====>.] - ETA: 0s - loss: 0.5639 - accuracy:
0.8419
    after epoch14,the accuracy is    0.84

Epoch 00015: val_accuracy did not improve from 0.82651
916/916 [=====] - 26s 28ms/step - loss: 0.5638 -
accuracy: 0.8419 - val_loss: 0.6683 - val_accuracy: 0.8242
Epoch 16/30
915/916 [=====>.] - ETA: 0s - loss: 0.5602 - accuracy:
0.8439
    after epoch15,the accuracy is    0.84

Epoch 00016: val_accuracy did not improve from 0.82651
916/916 [=====] - 25s 28ms/step - loss: 0.5603 -
accuracy: 0.8439 - val_loss: 0.6660 - val_accuracy: 0.8253
Epoch 17/30
915/916 [=====>.] - ETA: 0s - loss: 0.5545 - accuracy:
0.8448
    after epoch16,the accuracy is    0.84

Epoch 00017: val_accuracy did not improve from 0.82651
916/916 [=====] - 25s 28ms/step - loss: 0.5543 -
accuracy: 0.8449 - val_loss: 0.6637 - val_accuracy: 0.8227
Epoch 18/30
915/916 [=====>.] - ETA: 0s - loss: 0.5499 - accuracy:
0.8462
    after epoch17,the accuracy is    0.85

Epoch 00018: val_accuracy did not improve from 0.82651
916/916 [=====] - 25s 28ms/step - loss: 0.5498 -
accuracy: 0.8462 - val_loss: 0.6761 - val_accuracy: 0.8247
Epoch 19/30
915/916 [=====>.] - ETA: 0s - loss: 0.5450 - accuracy:
0.8466
    after epoch18,the accuracy is    0.85

Epoch 00019: val_accuracy improved from 0.82651 to 0.82658, saving model to
checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 27s 29ms/step - loss: 0.5450 -
accuracy: 0.8465 - val_loss: 0.6624 - val_accuracy: 0.8266
Epoch 20/30
915/916 [=====>.] - ETA: 0s - loss: 0.5411 - accuracy:
0.8473
    after epoch19,the accuracy is    0.85

Epoch 00020: val_accuracy did not improve from 0.82658

```


916/916 [=====] - 26s 28ms/step - loss: 0.5410 - accuracy: 0.8473 - val_loss: 0.6591 - val_accuracy: 0.8242
Epoch 21/30
915/916 [=====>.] - ETA: 0s - loss: 0.5397 - accuracy: 0.8495
after epoch20,the accuracy is 0.85

Epoch 00021: val_accuracy improved from 0.82658 to 0.82890, saving model to checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 26s 28ms/step - loss: 0.5396 - accuracy: 0.8495 - val_loss: 0.6617 - val_accuracy: 0.8289
Epoch 22/30
915/916 [=====>.] - ETA: 0s - loss: 0.5327 - accuracy: 0.8498
after epoch21,the accuracy is 0.85

Epoch 00022: val_accuracy improved from 0.82890 to 0.82924, saving model to checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 26s 28ms/step - loss: 0.5327 - accuracy: 0.8498 - val_loss: 0.6589 - val_accuracy: 0.8292
Epoch 23/30
914/916 [=====>.] - ETA: 0s - loss: 0.5287 - accuracy: 0.8500
after epoch22,the accuracy is 0.85

Epoch 00023: val_accuracy did not improve from 0.82924
916/916 [=====] - 26s 28ms/step - loss: 0.5284 - accuracy: 0.8500 - val_loss: 0.6889 - val_accuracy: 0.8155
Epoch 24/30
914/916 [=====>.] - ETA: 0s - loss: 0.5253 - accuracy: 0.8517
after epoch23,the accuracy is 0.85

Epoch 00024: val_accuracy improved from 0.82924 to 0.83019, saving model to checkpoint_best_only_cnn/checkpoint
916/916 [=====] - 26s 28ms/step - loss: 0.5258 - accuracy: 0.8515 - val_loss: 0.6538 - val_accuracy: 0.8302
Epoch 25/30
916/916 [=====] - ETA: 0s - loss: 0.5216 - accuracy: 0.8526
after epoch24,the accuracy is 0.85

Epoch 00025: val_accuracy did not improve from 0.83019
916/916 [=====] - 26s 28ms/step - loss: 0.5216 - accuracy: 0.8526 - val_loss: 0.6843 - val_accuracy: 0.8148
Epoch 26/30
915/916 [=====>.] - ETA: 0s - loss: 0.5171 - accuracy: 0.8532

after epoch25,the accuracy is 0.85

Epoch 00026: val_accuracy did not improve from 0.83019

916/916 [=====] - 26s 28ms/step - loss: 0.5172 - accuracy: 0.8532 - val_loss: 0.6598 - val_accuracy: 0.8245

Epoch 27/30

914/916 [=====>.] - ETA: 0s - loss: 0.5117 - accuracy: 0.8546

after epoch26,the accuracy is 0.85

Epoch 00027: val_accuracy did not improve from 0.83019

916/916 [=====] - 26s 28ms/step - loss: 0.5118 - accuracy: 0.8545 - val_loss: 0.6577 - val_accuracy: 0.8273

Epoch 28/30

915/916 [=====>.] - ETA: 0s - loss: 0.5081 - accuracy: 0.8548

after epoch27,the accuracy is 0.85

Epoch 00028: val_accuracy did not improve from 0.83019

916/916 [=====] - 26s 28ms/step - loss: 0.5084 - accuracy: 0.8547 - val_loss: 0.6503 - val_accuracy: 0.8294

Epoch 29/30

914/916 [=====>.] - ETA: 0s - loss: 0.5034 - accuracy: 0.8572

after epoch28,the accuracy is 0.86

Epoch 00029: val_accuracy did not improve from 0.83019

916/916 [=====] - 26s 28ms/step - loss: 0.5037 - accuracy: 0.8571 - val_loss: 0.6546 - val_accuracy: 0.8273

Epoch 30/30

914/916 [=====>.] - ETA: 0s - loss: 0.4995 - accuracy: 0.8578

after epoch29,the accuracy is 0.86

Epoch 00030: val_accuracy did not improve from 0.83019

916/916 [=====] - 26s 29ms/step - loss: 0.4998 - accuracy: 0.8578 - val_loss: 0.6571 - val_accuracy: 0.8294

```
[21]: fig_cnn=plt.figure(figsize=(12,5))

fig_cnn.add_subplot(121)

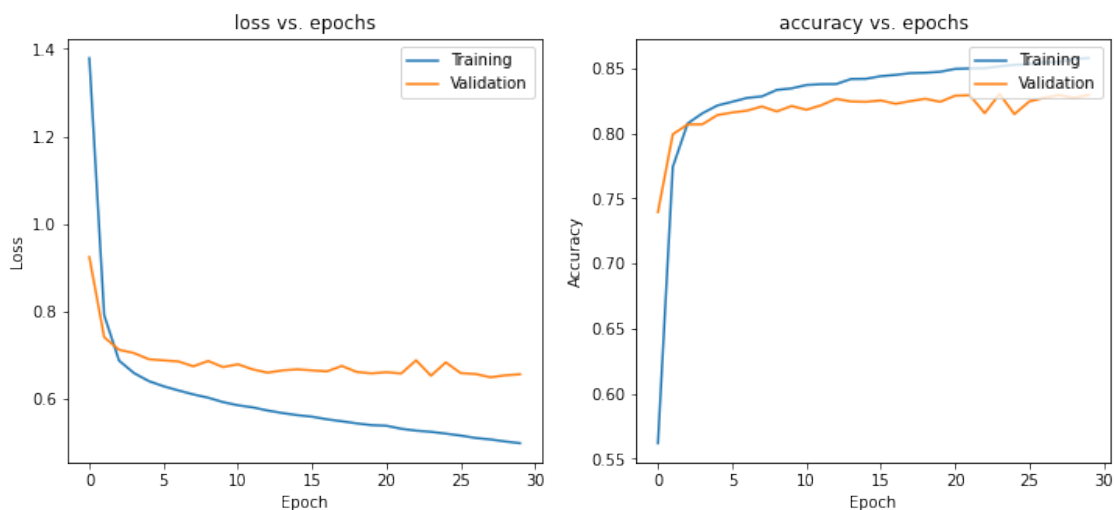
plt.plot(history_cnn.history['loss'])
plt.plot(history_cnn.history['val_loss'])
plt.title('loss vs. epochs')
plt.ylabel('Loss')
plt.xlabel('Epoch')
```

```
plt.legend(['Training', 'Validation'], loc='upper right')

fig_cnn.add_subplot(122)

plt.plot(history_cnn.history['accuracy'])
plt.plot(history_cnn.history['val_accuracy'])
plt.title('accuracy vs. epochs')
plt.ylabel('Accuracy')
plt.xlabel('Epoch')
plt.legend(['Training', 'Validation'], loc='upper right')

plt.show()
```



```
[22]: model_cnn.evaluate(test_data_gray, test_labels)
```

```
814/814 [=====] - 5s 7ms/step - loss: 0.7515 -
accuracy: 0.8085
```

```
[22]: [0.7514657974243164, 0.80854332447052]
```

1.5 4. Get model predictions

- Load the best weights for the MLP and CNN models that you saved during the training run.
- Randomly select 5 images and corresponding labels from the test set and display the images with their labels.
- Alongside the image and label, show each model's predictive distribution as a bar chart, and the final model prediction given by the label with maximum probability.


```
[23]: from tensorflow.keras.models import load_model
      from tensorflow.keras.preprocessing.image import img_to_array
      from tensorflow.keras.applications.resnet50 import preprocess_input,
      ↳ decode_predictions
```

```
[31]: model_mlp=get_new_model(train_data_gray[0].shape)
      model.load_weights('checkpoint_best_only/checkpoint')
      model_cnn2=Sequential([
          ↳
          ↳ Conv2D(16,kernel_size=3,activation='tanh',input_shape=(32,32,1)),
            MaxPooling2D(pool_size=3),
            Dense(64,activation='tanh'),
            Flatten(),
            Dense(10,activation='softmax')
      ])
      model_cnn2.load_weights('checkpoint_best_only_cnn/checkpoint')
```

```
[31]: <tensorflow.python.training.tracking.util.CheckpointLoadStatus at
      0x7f7178fd36d8>
```


```
[144]: np.random.seed(0)
      random_idx_2=np.random.choice(test_data_gray.shape[0],5)
      fig,axes=plt.subplots(1,5,figsize=(5,1))
      for i in range(5):
          axes[i].set_axis_off()
          axes[i].imshow(np.squeeze(test_data_gray[random_idx_2[i]]))
          axes[i].text(2., -1.5, f'Digit {numpy.
          ↳ argmax(test_labels[random_idx_2[i],])}')

```



```
[201]: def get_all_predictions(model):
      fig, ax = plt.subplots(1,5,figsize=(20,5))
      for i in range(5):
          prob=model.predict(img_to_array(test_data_gray[random_idx_2[i]])[np.
          ↳ newaxis, ...]).tolist()[0]
          ax[i].bar(np.insert(np.arange(9)+1,0,10),prob)
          ax[i].set_title(f'prediction label{np.argmax(prob)}')
```

```
[203]: get_all_predictions(model_mlp)
```



[204]: `get_all_predictions(model_cnn2)`

